

TOMAS FEITH

tsfeith@gmail.com $\diamond$ [linkedin.com/in/tomas-feith/](https://www.linkedin.com/in/tomas-feith/)
tomas-feith.github.io $\diamond$ github.com/tomas-feith

SKILLS

ML / Training: SFT | RLHF / RLAIF | Distillation | Synthetic Data | Evaluation

Infra / Inference: PyTorch | HuggingFace | TensorRT | Multi-Node GPU Training | Distributed Evaluation

Engineering: Python | C++ | TypeScript | SQL | Git | Docker | Linux

Languages: English (Fluent) | Portuguese (Native) | Spanish (Intermediate) | Dutch (Beginner)

EXPERIENCE

Senior Research Scientist

DeepL

Sep 2024 – Present
Rotterdam, Netherlands

- Led DeepL Write overhaul across all languages, shipping models that win the majority of head-to-head comparisons against every dedicated competitor and reach parity with GPT-4-class LLMs on internal preference evals
- Launched 3 new Write languages (JA/ZH/KO) end-to-end, unlocking 3 new key markets, by owning SFT, evaluation, and production rollout
- Built a multi-stage pipeline for PowerPoint translation, robustly translating heterogeneous slide layouts in production, by chaining extraction, translation, and slide reconstruction
- Shipped a document-translation evaluation framework, cutting eval turnaround from days to hours, replacing a manual qualitative process with a reliable, quantitative signal

Master's Thesis Researcher

EPFL, Data Science Lab (DLAB)

Sep 2023 – Feb 2024
Lausanne, Switzerland

- Introduced and formalised the *entity insertion* task in collaboration with Wikimedia, securing a first-author [EMNLP 2024](#) paper and an EPFL Best Thesis recommendation
- Built a 958M-link multilingual benchmark across 105 Wikipedia languages and open-sourced [code and dataset](#), by scraping the full Wikipedia with a custom filtering pipeline
- Trained xLocEI to nearly double GPT-4 (SotA at the time) on hits@1 (0.73 vs 0.37 macro-avg) at a fraction of the inference cost, by fine-tuning an XLM-RoBERTa-base encoder on the multilingual benchmark

Machine Learning Research Intern

Oracle Labs

Feb 2023 – Aug 2023
Zurich, Switzerland

- Filed 6 US patent applications in 6 months, by inventing novel joint pre-training objectives (AST/graph-path + next-token), multi-decoder classification, and cursor-aware doc re-ranking
- Pre-trained an efficient code-embedding LLM that cut inference latency 50%+ vs strong baselines, by curating a 10M+ source-code corpus from scratch and training with a novel joint objective
- Shipped a VSCode extension surfacing 1000+ docs pages in-IDE in real time, by integrating a VSCode front-end with a Flask back-end serving code embeddings from the pre-trained model

Student Researcher

EPFL, Multimedia Signal Processing Group

Sep 2022 – Feb 2023
Lausanne, Switzerland

- Designed a novel 3D-CNN architecture for light field denoising, securing publication at [SPIE 2023](#), by exploiting the camera grid across views to fuse multi-view information
- Rewrote the patch-matching core from Python to C++ with custom heuristics, achieving substantial end-to-end runtime improvement over the original implementation

EDUCATION

MSc of Computational Science and Engineering, EPFL, Switzerland

2024

BEng of Engineering Physics, Instituto Superior Técnico, Portugal

2021

PUBLICATIONS & PATENTS

6 US Patents filed (2023) on code-aware LLMs, joint pre-training, and ranking architectures. First-author papers: [Entity Insertion in Multilingual Linked Corpora](#) (*EMNLP 2024*); [Towards Learning-Based Denoising of Light Fields](#) (*SPIE 2023*).